

THE RESYNTHESIS CHANNEL: FUNDAMENTAL LIMITS OF SPEECH WATERMARKING AGAINST GENERATIVE LAUNDERING

Weiping Yan

University of Minnesota, Twin Cities, Minneapolis, MN, USA

ABSTRACT

Audio watermarks are increasingly defeated when an attacker regenerates the signal through a vocoder or codec; this “regeneration” collapse is widely documented, yet for speech there is no theory of *which* watermarks die, under *which* attacker, or *what can survive*. We supply one. We model any blind analysis–resynthesis launderer as a *resynthesis channel* $\mathcal{W} = \mathcal{S} \circ \mathcal{A}$: an analysis map $\mathcal{A}$ that keeps only the perceptual invariants of speech (content, speaker, prosody) and a synthesis map $\mathcal{S}$ that regenerates a waveform, re-sampling everything $\mathcal{A}$ discarded. Two matched results follow. A *converse* (data-processing inequality): any watermark supported in $\ker \mathcal{A}$ —phase and psychoacoustically masked components, where imperceptibility drives post-hoc schemes—has post-laundering detection error-exponent exactly zero. An *achievability*: a watermark aligned to the preserved invariants survives at rate R^* , the capacity of the invariant sub-channel. The survivable set is exactly the non-nullspace of $\mathcal{A}$. On LibriSpeech, a surface mark and AudioSeal collapse under lossy resynthesis (mel-spectrogram inversion, low-rate EnCodec) while an invariant-aligned mark survives and degrades gracefully as the channel keeps less; post-laundering detectability is a single monotone function of the invariant-energy fraction, as the converse predicts.

Index Terms— audio watermarking, generative AI, provenance, data-processing inequality, hypothesis testing

1. INTRODUCTION

Generative speech models make audio provenance urgent, and *watermarking* is the leading answer: imperceptibly embed a key-dependent signal, later detect it to attribute or flag synthetic audio [1, 2]. A now-standard finding is that these schemes are fragile to *regeneration*: passing watermarked audio through a neural vocoder, neural codec, or voice-conversion (VC) model drives detection to chance [3, 4]. This is not incidental noise—it is the natural operation of any generative model, and it defeats watermarks that survive conventional distortions (noise, MP3, resampling).

The phenomenon is thoroughly *measured* but, for audio, not *explained*. For images this is well understood: Zhao *et al.* [5] prove regeneration removes pixel-space watermarks and prescribe embedding in preserved *semantic* content, and semantic/latent-space image watermarks (Tree-Ring [6], Gaussian Shading [7]) realize exactly that. What is missing—and what we provide—is a *quantitative theory for speech*: a matched converse and achievability that pin down not just *that* surface marks die but *how much* detectability survives and at what *rate* R^* , exploiting the concrete, estimable invariants of audio analysis–resynthesis (content, speaker, prosody) and the phase/masking geometry specific to hearing.

Contributions. (i) We cast blind generative laundering as a *resynthesis channel* $\mathcal{W} = \mathcal{S} \circ \mathcal{A}$ (§2). (ii) A *converse* (Thm. 1): watermark energy in $\ker \mathcal{A}$ has zero surviving detection exponent—a data-processing inequality that explains the collapse, sharpened for audio by phase/masking. (iii) A matched *achievability* (Thm. 2): an invariant-aligned watermark survives at rate R^* , the invariant sub-channel capacity; the survivable set is exactly the non-nullspace of $\mathcal{A}$. (iv) Validation on real speech (§5): across analysis–resynthesis channels of increasing nullspace (STFT and mel-spectrogram inversion, and an EnCodec codec), a surface mark and AudioSeal collapse when their energy lies in the channel’s nullspace, while an invariant mark survives and degrades gracefully as the preserved subspace shrinks; post-laundering detectability is monotone in the invariant-energy fraction, as the converse predicts.

2. THE RESYNTHESIS CHANNEL

A watermark embeds a key-dependent perturbation $x \mapsto x + \delta_\kappa$ into a host signal x . A detector that knows the key κ decides H_0 (clean, law P_0) vs. H_1 (watermarked, law P_1). The optimal Bayesian error after N i.i.d. looks decays as $e^{-N C(P_0, P_1)}$ with the *Chernoff information* $C(P_0, P_1) = \max_{s \in [0, 1]} -\log \int p_0^{1-s} p_1^s$ [8, 9]; $C = 0$ iff $P_0 = P_1$, i.e. no detector beats chance at any N .

Attacker. A blind launderer applies $\mathcal{W} = \mathcal{S} \circ \mathcal{A}$: analysis $\mathcal{A} : x \mapsto z$ extracts perceptual invariants z (a near-sufficient statistic for perceptual identity—two signals with equal z sound the same), and synthesis $\mathcal{S} : z \mapsto \hat{x}$ regenerates a waveform consistent with z , drawing everything $\mathcal{A}$ discarded from the clean prior $P_0(\cdot | z)$. The attacker does not know κ and does not optimize against the scheme (§6). Imperceptibility is a budget $\rho(\delta) \leq D$ for a psychoacoustic distortion ρ .

Linear–Gaussian surrogate. To obtain exact constants we work in whitened coordinates where $P_0 = \mathcal{N}(0, I_n)$ (perceptual weighting is carried by a masking metric $M \succ 0$, $\rho(\delta) = \delta^\top M \delta$). Let $\mathcal{A} \in \mathbb{R}^{k \times n}$ have full row rank, $P = \mathcal{A}^+ \mathcal{A}$ the orthogonal projector onto $\text{row}(\mathcal{A})$ (the *invariant subspace*), and $I - P$ the projector onto $\ker \mathcal{A}$ (the *surface subspace* $\mathcal{S}$ resamples). Then

$$\mathcal{W}(x) = Px + (I - P)w, \quad w \sim \mathcal{N}(0, I_n) \text{ fresh}, \quad (1)$$

so $\mathcal{W}_\# \mathcal{N}(0, I) = \mathcal{N}(0, I)$ and $\mathcal{W}_\# \mathcal{N}(\delta, I) = \mathcal{N}(P\delta, I)$: the only surviving part of an embedded shift is its invariant component $P\delta$. (The first clause of Thm. 1 needs no Gaussianity—Chernoff information obeys the DPI for any channel; the surrogate supplies constants and the achievable construction.)

3. CONVERSE: WHAT IS ERASED

Theorem 1 (Nullspace erasure). *For the channel (1) and any watermark δ , $C(\mathcal{W}_\# P_0, \mathcal{W}_\# P_1) = \frac{1}{8} \|P\delta\|^2 \leq \frac{1}{8} \|\delta\|^2 = C(P_0, P_1)$,*

Code, proofs, and data: github.com/appleweiping/resynthesis-watermark-limits.

with equality iff $\delta \in \text{row}(\mathcal{A})$. In particular, if $\delta \in \ker \mathcal{A}$ then $C(\mathcal{W}_\# P_0, \mathcal{W}_\# P_1) = 0$: after laundering no detector beats chance, for any N . Moreover $\mathcal{W}(x + \delta) \stackrel{d}{=} \mathcal{W}(x)$, so the laundered watermarked signal is distributed identically to a laundered clean one.

Proof sketch. $\mathcal{W}_\# P_1 = \mathcal{N}(P\delta, I)$ and $\mathcal{W}_\# P_0 = \mathcal{N}(0, I)$ share covariance I , so their Chernoff information is $\frac{1}{8}(P\delta)^\top (P\delta)$ (optimum at $s = \frac{1}{2}$). P is an orthogonal projector, hence $\|P\delta\| \leq \|\delta\|$, giving the DPI bound, tight on $\text{row}(\mathcal{A})$. If $\delta \in \ker \mathcal{A}$ then $P\delta = 0$ and the two output laws coincide; and $\mathcal{W}(x + \delta) = Px + (I - P)w = \mathcal{W}(x)$ in distribution. $\square$

Remark 1 (Why the nullspace is huge for audio). *Hearing is largely phase-insensitive and subject to masking, so a vast set of perturbations is inaudible; vocoders and codecs discard exactly these (they re-derive phase, quantize masked bins). Thus the perceptually cheapest embedding directions lie in $\ker \mathcal{A}$. Post-hoc watermarks, optimized for imperceptibility, drift there—and die. This is the theory’s account of the empirical collapse [3, 4].*

4. ACHIEVABILITY: WHAT SURVIVES

Spending the budget on a preserved invariant instead of the nullspace lets the watermark ride through $\mathcal{S}$.

Theorem 2 (Invariant sub-channel). *Under the budget, the best surviving exponent is $\max_{\delta^\top M \delta \leq D} \frac{1}{8} \|P\delta\|^2 = \frac{D}{8} \lambda_{\max}(P, M)$, at the top generalized eigenvector of (P, M) ; it is 0 when $\delta \in \ker \mathcal{A}$. For payload, encoding a message as a surviving shift $u = P\delta \in \text{row}(\mathcal{A})$ and facing the clean host as interference ($y = u + \eta$, $\eta \sim \mathcal{N}(0, I)$) gives surviving rate*

$$R^* = \max_{\text{tr}(M_{\text{row}} Q) \leq D} \frac{1}{2} \log \det(I + Q), \quad (2)$$

a water-filling over the eigenmodes of $M_{\text{row}} = V_{\text{row}}^\top M V_{\text{row}}$. $R^* > 0$ whenever $D > 0$; a nullspace watermark has $R^* = 0$ after $\mathcal{W}$.

Theorem 1 caps the survivable set at $\text{row}(\mathcal{A})$; Theorem 2 achieves a positive rate on it. The survivable set is exactly the non-nullspace of $\mathcal{A}$, of size R^* .

5. EXPERIMENTS

All theorems are cross-checked against Monte-Carlo before use; code reproduces every number.

Surrogate (validation). Fig. 1 instantiates (1) with random invariant subspaces. As a watermark drifts from the invariant subspace into $\ker \mathcal{A}$, post-laundering AUC falls from >0.88 to 0.50 (chance), matching the closed form to RMSE 0.001; the surviving detection exponent and rate R^* are positive for the invariant-aligned mark and identically zero for the post-hoc (nullspace) mark.

Real speech. We use 80 LibriSpeech test-clean utterances [10] at 16 kHz. Attackers instantiate $\mathcal{W}$ with *increasing nullspace*: a full linear-STFT Griffin-Lim [11] (a near-lossless control: $\ker \mathcal{A}$ is only the phase), a mel-spectrogram Griffin-Lim (lossy: $\ker \mathcal{A}$ is phase *plus* all sub-mel detail), and EnCodec [12] round-trips at 6/3/1.5 kbps (nullspace grows as the rate drops). Watermarks at matched SNR (24 dB): a *surface* mark (fine high-band carrier, waveform-correlation detector, in $\ker \mathcal{A}$), an *invariant* mel-envelope mark (the constructive realization of Thm. 2), and *AudioSeal*, a deployed neural watermark [1]. Each mark’s invariant-energy fraction




Fig. 1. Surrogate. Left: surviving Chernoff exponent vs. imperceptibility budget for the invariant-aligned vs. post-hoc (nullspace) watermark. Right: surviving payload rate R^* ; the nullspace rate is identically zero.

Table 1. Real speech (LibriSpeech, $N = 80$). Detection AUC before $\rightarrow$ after each blind attacker. STFT[†] is a near-lossless control (nullspace = phase only); mel-GL and EnCodec (EnCk, k kbps) are lossy. Surface/AudioSeal (energy in the nullspace) collapse to chance; the invariant mel mark survives.

Watermark	STFT [†]	mel-GL	EnC6	EnC3	EnC1.5
Surface (null)	1.00 $\rightarrow$ 0.48	1.00 $\rightarrow$ 0.44	1.00 $\rightarrow$ 0.53	1.00 $\rightarrow$ 0.54	1.00 $\rightarrow$ 0.56
Invariant (mel)	0.94 $\rightarrow$ 0.93	0.94 $\rightarrow$ 0.93	0.94 $\rightarrow$ 0.74	0.94 $\rightarrow$ 0.67	0.94 $\rightarrow$ 0.60
AudioSeal	1.00 $\rightarrow$ 1.00	1.00 $\rightarrow$ 0.20	1.00 $\rightarrow$ 0.99	1.00 $\rightarrow$ 0.98	1.00 $\rightarrow$ 0.94

f (STFT-magnitude change captured by the mel envelope) is a pure signal-geometry analog of $\|P\delta\|^2 / \|\delta\|^2$.

Table 1 gives Part A, and it sweeps exactly the nullspace-size axis. The STFT control has the *smallest* nullspace (phase only): it already erases the phase-based surface mark (AUC $\rightarrow$ 0.5) yet preserves both magnitude-domain marks. Enlarging the nullspace to phase *plus* sub-mel detail (mel-spectrogram inversion) additionally collapses AudioSeal (AUC 1.00 $\rightarrow$ 0.203, bit-accuracy 0.482; the sub-0.5 AUC is a detector sign-flip, detection destroyed either way), while the invariant mel mark is untouched (0.936 $\rightarrow$ 0.927). Shrinking the invariant subspace further with low-rate EnCodec then *gracefully erodes* even the invariant mark (0.74/0.67/0.60 at 6/3/1.5 kbps)—the behaviour of R^* in Thm. 2 as the channel keeps less. Here f is measured against the *mel* analysis map, so it predicts the mel-inversion outcomes; each attacker has its own preserved subspace (a channel-relative $f_{\mathcal{W}}$), which is why AudioSeal ($f = 0.262$, little energy in the mel envelope) resists the EnCodec it was hardened against yet dies under mel-inversion, whose nullspace contains the phase and fine detail it relies on. The invariant mark ($f = 0.812$) sits in the mel-preserved subspace. Part B makes this quantitative: sweeping a mixture from surface to invariant, post-laundering AUC is a monotone function of f , well described by the converse’s $\sqrt{f}$ form $\Phi(a\sqrt{f} + b)$ (Fig. 2)—near chance until f exceeds the surface floor, then rising as the preserved invariant energy grows, as Thm. 1 predicts.

Achievability & rate. Carrying a multi-bit payload in the same invariant sub-channel (BPSK over mel-envelope patterns), the invariant mark recovers bits after laundering where a surface payload of equal size decays to chance; the surviving payload grows with the imperceptibility budget and shrinks with the attacker’s aggressiveness, the qualitative signature of the water-filling rate R^* (2) (Fig. 3).




Fig. 2. Real speech (mel-inversion channel). Post-laundering detection AUC vs. the invariant-energy fraction f . A surface→invariant mixture (circles) and the named marks (triangles) follow the converse’s $\sqrt{f}$ form $\Phi(a\sqrt{f} + b)$ (two-parameter fit); low- f marks collapse to chance, the high- f invariant mark survives.

6. DISCUSSION AND SCOPE

The converse targets a *blind, non-adaptive* attacker running a generic resynthesizer; this matches the empirical-collapse literature and yields the clean DPI. A key-aware attacker that optimizes to remove our specific mark is an arms race we do not claim to win—we treat it as a stress test, not a core result. Real analysis maps are only approximately sufficient and $\mathcal{S}$ is stochastic; our claims degrade gracefully as these idealizations break. Two honest limitations: marks are matched by SNR rather than a perceptual metric (PESQ/ViSQOL), so part of the invariant mark’s survival may be bought with audibility; and our synthesizers are spectral-inversion and a neural codec—a learned neural vocoder and voice conversion would test the converse in its hardest regime. The practical message is prescriptive: *a watermark survives generative laundering only insofar as it perturbs the invariants the generator is built to preserve*—the regime where imperceptibility is hardest, quantified by R^* .

7. RELATION TO PRIOR WORK

Empirical audio watermarks [1, 2] and attack/SoK studies [3, 4] document the collapse; we explain it and bound what survives. In the image domain the story is understood: Zhao *et al.* [5] prove removability and advocate semantic embedding, and Tree-Ring [6] and Gaussian Shading [7] realize regeneration-robust semantic watermarks. Our contribution is not the qualitative “embed-in-what-survives” idea but a *quantitative, provable* matched converse and achievability with an explicit surviving rate R^* for the *speech* resynthesis channel, tied to its concrete, measurable invariants. Classical information-theoretic watermarking [13, 14, 15] treats fixed additive attacks with informed embedding; our attacker is a nonlinear manifold projection and our detector is blind, locating survivability in the geometry of $\mathcal{A}$.

8. CONCLUSION

Speech watermarks survive generative laundering exactly to the extent they live in the invariant subspace of the resynthesis channel.




Fig. 3. Achievability. Payload bit-accuracy after mel-inversion laundering vs. imperceptibility budget (SNR). The invariant sub-channel carries a surviving payload; a surface payload of equal size decays to chance.

We proved a matched converse and achievability and confirmed on spectral-inversion and neural-codec channels that detection collapses as the surviving energy vanishes. Designing watermarks *in* the invariant sub-channel, at rate R^* , is the path to laundering-robust provenance.

9. REFERENCES

- [1] Robin San Roman, Pierre Fernandez, Alexandre Défossez, Teddy Furon, Tuan Tran, and Hady Elsahar, “Proactive detection of voice cloning with localized watermarking,” in *International Conference on Machine Learning (ICML)*, 2024.
- [2] Guangyu Chen, Yu Wu, Shujie Liu, Tao Liu, Xiaoyong Du, and Furu Wei, “WavMark: Watermarking for audio generation,” 2023, arXiv:2308.12770.
- [3] Patrick O’Reilly, Zeyu Jin, Jiaqi Su, and Bryan Pardo, “Deep audio watermarks are shallow: Limitations of post-hoc watermarking techniques for speech,” in *ICLR Workshop on GenAI Watermarking*, 2025.
- [4] Yizhu Wen, Ashwin Innuganti, Aaron Bien Ramos, Hanqing Guo, and Qiben Yan, “SoK: How robust is audio watermarking in generative AI models?,” 2025, arXiv:2503.19176.
- [5] Xuandong Zhao, Kexun Zhang, Zihao Su, Saastha Vasani, Ilya Grishchenko, Christopher Kruegel, Giovanni Vigna, Yu-Xiang Wang, and Lei Li, “Invisible image watermarks are provably removable using generative AI,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.
- [6] Yuxin Wen, John Kirchenbauer, Jonas Geiping, and Tom Goldstein, “Tree-ring watermarks: Fingerprints for diffusion images that are invisible and robust,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- [7] Zijin Yang, Kai Zeng, Kejiang Chen, Han Fang, Weiming Zhang, and Nenghai Yu, “Gaussian Shading: Provable performance-lossless image watermarking for diffusion models,” in *IEEE/CVF Conf. on Computer Vision and Pattern Recognition (CVPR)*, 2024.
- [8] Herman Chernoff, “A measure of asymptotic efficiency for tests of a hypothesis based on the sum of observations,” *The Annals of Mathematical Statistics*, vol. 23, no. 4, pp. 493–507, 1952.
- [9] Thomas M. Cover and Joy A. Thomas, *Elements of Information Theory*, Wiley-Interscience, 2nd edition, 2006.
- [10] Vassil Panayotov, Guoguo Chen, Daniel Povey, and Sanjeev Khudanpur, “Librispeech: An ASR corpus based on public domain audio books,” in *IEEE Int. Conf. on Acoustics, Speech and Signal Processing (ICASSP)*, 2015.
- [11] Daniel Griffin and Jae Lim, “Signal estimation from modified short-time Fourier transform,” *IEEE Transactions on Acoustics, Speech, and Signal Processing*, vol. 32, no. 2, pp. 236–243, 1984.
- [12] Alexandre Défossez, Jade Copet, Gabriel Synnaeve, and Yossi Adi, “High fidelity neural audio compression,” *Transactions on Machine Learning Research (TMLR)*, 2023.
- [13] Pierre Moulin and Joseph A. O’Sullivan, “Information-theoretic analysis of information hiding,” *IEEE Transactions on Information Theory*, vol. 49, no. 3, pp. 563–593, 2003.
- [14] Max Costa, “Writing on dirty paper,” *IEEE Transactions on Information Theory*, vol. 29, no. 3, pp. 439–441, 1983.
- [15] Brian Chen and Gregory W. Wornell, “Quantization index modulation: A class of provably good methods for digital watermarking and information embedding,” *IEEE Transactions on Information Theory*, vol. 47, no. 4, pp. 1423–1443, 2001.